

VIETNAM NATIONAL UNIVERSITY - HO CHI MINH CITY
HO CHI MINH CITY UNIVERSITY OF TECHNOLOGY
FACULTY OF COMPUTER SCIENCE AND ENGINEERING



PROBABILITY AND STATISTICS (MT2013) - SEMESTER 231

Project Topic

Computer Parts (CPUs and GPUs)

Lecturer: Nguyen Tien Dung
Class: CC02 Group: 6
Students: Nong Thuc Khue - 2252385
Vu Quynh Huong - 2252286
Le Phong Hao - 2252182
Huynh Van Anh Hoang - 2252229
Truong Minh Khang - 2252314

HO CHI MINH CITY, DECEMBER 2023



Contribution

No.	Name	Student ID	Task	Completion
1	Nong Thuc Khue	2252385	Code writer	100%
2	Vu Quynh Huong	2252286	Code writer	100%
3	Le Phong Hao	2252182	Report management	100%
4	Huynh Van Anh Hoang	2252229	Report management	100%
5	Truong Minh Khang	2252314	Report management	100%



List of Figures

1	Histograms of data features	18
2	Scatter plots of data features	20
3	Correlation matrix (1)	21
4	Correlation matrix (2)	21
5	Correlation matrix (3)	22
6	Correlation matrix (4)	22
7	Correlation matrix (5)	23

Contents

1	Theoretical foundation	4
1.1	Linear regression	4
1.1.1	Two types of linear regression	4
1.1.2	Linear regression model	4
1.1.3	Cost function	4
1.1.4	Gradient descent	5
1.1.5	Model performance	5
1.2	Stepwise regression	5
1.2.1	Definition of stepwise regression	5
1.2.2	Stepwise regression and Linear regression	5
1.3	Analysis of variance (ANOVA)	6
1.3.1	Definition of ANOVA	6
1.3.2	What does the Analysis of Variance (ANOVA) reveal?	6
1.3.3	One-way ANOVA versus Two-way ANOVA	7
1.4	Support vector machine	7
1.4.1	Definition of support vector machine	7
1.4.2	Types of support vector machine	8
2	Data processing	9
2.1	Data importing	9
2.1.1	Import file	9
2.1.2	Import libraries	10
2.1.3	Explanation of the dataset	10
2.2	Data cleaning	10
2.2.1	Removing unused data	10
2.2.2	Filling missing values	11
2.3	Data visualization	15
2.3.1	Descriptive statistics for each variable	15
2.3.2	Plotting graph	16
2.3.2.a	Histogram	16
2.3.2.b	Box plot	18
2.3.2.c	Correlation matrix	20
2.4	Models building	23
2.4.1	Multivariate Linear Regression (MLR)	23
2.4.2	Analysis of Variance (ANOVA)	25
2.4.3	Support vector machine (SVM)	26
2.5	Model comparison	26
3	Conclusion	28
4	R script	29

1 Theoretical foundation

1.1 Linear regression

1.1.1 Two types of linear regression

Single Linear Regression: If a single independent variable is used to predict the value of a numerical dependent variable, then such a Linear Regression algorithm is called Simple Linear Regression.

The expected value of Y at each level of x is a random variable $\mathbb{E}(Y|x) = \alpha + \beta x$. We assume that each observation, Y , can be described by the model $Y = \alpha + \beta x + \epsilon$. That is, $Y_i = \alpha + \beta x_i + \epsilon_i, i = 1, 2, 3, \dots, n$.

To achieve the best-fit regression line, it is essential to make predictions the value that the difference between the predicted value and the true value Y is the *minimum*. So, we have to update the value of α and β , to reach the best values of both ones.

Multiple Linear Regression: If more than one independent variable is used to predict the value of a numerically dependent variable, then such a Linear Regression algorithm is called Multiple Linear Regression.

1.1.2 Linear regression model

Linear regression is a powerful tool for understanding and predicting the behavior of a variable, however, it needs to meet a few conditions in order to be accurate and dependable solutions.

- **Linearity:** The independent and dependent variables have a linear relationship with each other. This implies that changes in the dependent variable follow those in the independent variable(s) in a linear fashion.
- **Independence:** The observations in the dataset are independent of each other. This means that the value of the dependent variable for one observation does not depend on the value of the dependent variable for another observation.
- **Homoscedasticity:** Across all levels of the independent variable(s), the variance of the errors is constant. This indicates that the amount of independent variable(s) does not have an impact on the variance of the errors.
- **Normality:** The errors in the model are normally distributed.
- **No multi-collinearity:** There is no high correlation between the independent variables. This indicates that there is little or no correlation between the independent variables.

1.1.3 Cost function

The cost function or the loss function is nothing more than the error or difference between the predicted \hat{Y} and the true value Y . It is the Mean Squared Error (MSE) between the predicted value and the true value. The joint function (J) can be written as:

$$\text{Cost function}(J) = \frac{1}{n} \sum_{i=1}^n (\hat{y}_i - y_i)^2$$

Residuals: The distance between the actual value and predicted values is called *residual*. If the observed points are far from the regression line, then the residual will be high, and so the cost function will be high. If the scatter points are close to the regression line, then the residual will be small and hence the cost function.

1.1.4 Gradient descent

Gradient descent is used to minimize the MSE by calculating the gradient of the cost function.

A regression model uses gradient descent to update the coefficients of the line by reducing the cost function. It is done by a random selection of values of coefficient and then iteratively update the values to reach the minimum cost function.

1.1.5 Model performance

The goodness of fit determines how the line of regression fits the set of observations. The process of finding the best model out of the various models is called optimization. It can be achieved by R-squared method:

- R-squared is a statistical method that determines the goodness of fit.
- It measures the strength of the relationship between the dependent and independent variables on a scale of 0 – 100%.
- The high value of R-square determines the less difference between the predicted values and actual values and hence represents a good model.
- It is also called a coefficient of determination, or coefficient of multiple determination for multiple regression.
- It can be calculated from the below formula:

$$R\text{-squared} = \frac{\text{Explained variation}}{\text{Total variation}}$$

1.2 Stepwise regression

1.2.1 Definition of stepwise regression

Stepwise regression is a method of fitting a regression model by iteratively adding or removing variables. It is used to build a model that is accurate and parsimonious, which means that it has the smallest number of variables that can explain the data. There are two main types of stepwise regression:

- **Forward Selection:** the algorithm starts with an empty model and iteratively adds variables to the model until no further improvement is made.
- **Backward Elimination:** the algorithm starts with a model that includes all variables and iteratively removes variables until no further improvement is made.

The advantage of stepwise regression is that it can automatically select the most important variables for the model and build a parsimonious model. The disadvantage is that it may not always select the best model and can be sensitive to the order in which the variables are added or removed.

1.2.2 Stepwise regression and Linear regression

Linear regression is a statistical method used to model the relationship between a dependent variable and one or more independent variables by fitting a linear equation to observed data. In other words, it is a method for predicting a response (or dependent variable) based on one or more predictor variables.

Stepwise regression is a method for building a regression model by adding or removing predictors in a step-by-step fashion. The goal of stepwise regression is to identify the subset of predictors that

provides the best predictive performance for the response variable. This is done by starting with an empty model and iteratively adding or removing predictors based on the strength of their relationship with the response variable.

In summary, linear regression is a method for modeling the relationship between a response and one or more predictor variables, while stepwise regression is a method for building a regression model by iteratively adding or removing predictors.

1.3 Analysis of variance (ANOVA)

1.3.1 Definition of ANOVA

Analysis of variance (ANOVA) is an analysis tool used in statistics that divides an observed aggregate variability found inside a data set into two parts: systematic factors and random factors. Systematic factors have a statistical influence on the given data set, while random factors do not.

Analysts use the ANOVA test to determine the influence that independent variables have on the dependent variable in a regression study.

The t- and z-test methods developed in the 20th century were used for statistical analysis until 1918, when Ronald Fisher created the analysis of variance method.

ANOVA is also called Fisher analysis of variance and is the extension of the t- and z-tests.

What are t-test and z-test?

Both serve as hypothesis tests that evaluate whether there is a notable difference between the means of two distinct groups or populations.

The t-test is used when the population variance is unknown, or the sample size is small ($n < 30$). At the same time, the z-test is applied when the population variance is known, and the sample size is large ($n > 30$).

T-test employs the Student's t-distribution, while z-test uses the standard normal distribution. As the sample size increases, t-distribution will converge into the standard normal distribution.

The t-test can be understood as a statistical test which is used to compare and analyze whether the means of the two population is different from one another or not when the standard deviation is not known. As against, z-test is a parametric test, which is applied when the standard deviation is known, to determine whether the means of the two datasets differ from each other.

1.3.2 What does the Analysis of Variance (ANOVA) reveal?

The ANOVA test is the initial step in analyzing the factors that affect a given data set. Once the test is finished, an analyst performs additional testing on the methodical factors that measurably contribute to the data set's inconsistency. The analyst utilizes the results of the ANOVA test in an f-test to generate additional data that align with the proposed regression models.

The ANOVA test allows a comparison of more than two groups at the same time to determine whether a relationship exists between them. The result of the ANOVA formula, the F statistic (also called the F-ratio), allows for the analysis of multiple groups of data to determine the variability between samples and within samples.

$$F = \frac{MSTr}{MSE} = \frac{\frac{SSTr}{I-1}}{\frac{SSE}{N-1}}$$

Where,

- $MSTr = \frac{SST}{I-1}$: mean square for treatment.

- $MSE = \frac{SSE}{N - 1}$: mean square for error.

If there is no real difference between the tested groups, which is called the null hypothesis, the result of the ANOVA's F-ratio statistic will be close to 1. The distribution of all possible values of the F statistic is the F-distribution. This is actually a group of distribution functions, with two characteristic numbers called the numerator degrees of freedom and the denominator degrees of freedom.

1.3.3 One-way ANOVA versus Two-way ANOVA

There are two main types of ANOVA: one-way (or unidirectional) and two-way. There also variations of ANOVA. For example, MANOVA (multivariate ANOVA) differs from ANOVA as the former tests for multiple dependent variables simultaneously while the latter assesses only one dependent variable at a time.

One-way or two-way refers to the number of independent variables in your analysis of variance test. A one-way ANOVA evaluates the impact of a sole factor on a sole response variable. It determines whether all the samples are the same.

One-way analysis of variance (ANOVA) is used to determine whether there are statistically significant differences between the means of three or more independent (unrelated) groups.

A two-way ANOVA is an extension of the one-way ANOVA. With a one-way, you have one independent variable affecting a dependent variable. With a two-way ANOVA, there are two independents. It is utilized to observe the interaction between the two factors and tests the effect of two factors at the same time.

Furthermore, ANOVA relies on assumptions. ANOVA tests assume that the data are normally distributed and that the levels of variance in each group are roughly equal. Finally, we assume that all observations are made independently. If these assumptions are not accurate, ANOVA may not be useful for comparing groups.

1.4 Support vector machine

1.4.1 Definition of support vector machine

Support Vector Machine (SVM) is a powerful machine learning algorithm used for linear or nonlinear classification, regression, and even outlier detection tasks. SVMs can be used for a variety of tasks, such as text classification, image classification, spam detection, handwriting identification, gene expression analysis, face detection, and anomaly detection. SVMs are adaptable and efficient in a variety of applications because they can manage high-dimensional data and nonlinear relationships.

SVM algorithms are very effective as we try to find the maximum separating hyperplane between the different classes available in the target feature.

Support Vector Machine (SVM) is a supervised machine learning algorithm used for both classification and regression. Although we also say regression problems, it is best suited for classification. The main objective of the SVM algorithm is to find the optimal hyperplane in an N-dimensional space that can separate the data points in different classes in the feature space. The hyperplane tries to make the margin between the closest points of different classes as maximum as possible. The dimension of the hyperplane depends on the number of features. If the number of input features is two, then the hyperplane is just a line. If the number of input features is three, then the hyperplane becomes a 2-D plane. It becomes difficult to imagine when the number of features exceeds three.

1.4.2 Types of support vector machine

Based on the nature of the decision boundary, Support Vector Machines (SVM) can be divided into two main parts:

- **Linear SVM:** Linear SVMs use a linear decision boundary to separate the data points of different classes. When the data can be precisely linearly separated, linear SVMs are very suitable. This means that a single straight line (in 2D) or a hyperplane (in higher dimensions) can entirely divide the data points into their respective classes. A hyperplane that maximizes the margin between the classes is the decision boundary.
- **Non-Linear SVM:** Non-Linear SVM can be used to classify data when it cannot be separated into two classes by a straight line (in the case of 2D). By using kernel functions, nonlinear SVMs can handle nonlinearly separable data. The original input data is transformed by these kernel functions into a higher-dimensional feature space, where the data points can be linearly separated. A linear SVM is used to locate a non-linear decision boundary in this modified space.

2 Data processing

2.1 Data importing

2.1.1 Import file

```
1 #read data
2 CPU_data<-read.csv("~/Desktop/Assignment/work/Intel_CPUs.csv")
3 head(CPU_data, 3)
```

	Product_Collection			Vertical_Segment	Processor_Number	
1	7th Generation Intel	Core i7	Processors		Mobile	i7-7Y75
2	8th Generation Intel	Core i5	Processors		Mobile	i5-8250U
3	8th Generation Intel	Core i7	Processors		Mobile	i7-8550U
	Status	Launch_Date	Lithography	Recommended_Customer_Price	nb_of_Cores	
1	Launched	Q3'16	14 nm	\$393.00	2	
2	Launched	Q3'17	14 nm	\$297.00	4	
3	Launched	Q3'17	14 nm	\$409.00	4	
	nb_of_Threads	Processor_Base_Frequency		Max_Turbo_Frequency	Cache	
1	4	1.30 GHz		3.60 GHz	4 MB	SmartCache
2	8	1.60 GHz		3.40 GHz	6 MB	SmartCache
3	8	1.80 GHz		4.00 GHz	8 MB	SmartCache
	Bus_Speed	TDP	Embedded_Options_Available	Conflict_Free	Max_Memory_Size	
1	4 GT/s OPI	4.5 W	No	Yes	16 GB	
2	4 GT/s OPI	15 W	No	Yes	32 GB	
3	4 GT/s OPI	15 W	No	Yes	32 GB	
	Memory_Types		Max_nb_of_Memory_Channels	Max_Memory_Bandwidth		
1	LPDDR3-1866, DDR3L-1600		2	29.8 GB/s		
2	DDR4-2400, LPDDR3-2133		2	34.1 GB/s		
3	DDR4-2400, LPDDR3-2133		2	34.1 GB/s		
	ECC_Memory_Supported		Processor_Graphics_	Graphics_Base_Frequency		
1	No		NA	300 MHz		
2	No		NA	300 MHz		
3	No		NA	300 MHz		
	Graphics_Max_Dynamic_Frequency		Graphics_Video_Max_Memory	Graphics_Output		
1	1.05 GHz		16 GB	eDP/DP/HDMI/DVI		
2	1.10 GHz		32 GB	eDP/DP/HDMI/DVI		
3	1.15 GHz		32 GB	eDP/DP/HDMI/DVI		
	Support_4k	Max_Resolution_HDMI	Max_Resolution_DP			
1	NA	4096x2304@24Hz	3840x2160@60Hz			
2	NA	4096x2304@24Hz	4096x2304@60Hz			
3	NA	4096x2304@24Hz	4096x2304@60Hz			
	Max_Resolution_eDP_Integrated_Flat_Panel		DirectX_Support	OpenGL_Support		
1	3840x2160@60Hz		12	NA		
2	4096x2304@60Hz		12	NA		
3	4096x2304@60Hz		12	NA		
	PCI_Express_Revision	PCI_Express_Configurations_		Max_nb_of_PCI_Express_Lanes		
1	3	1x4, 2x2, 1x2+2x1 and 4x1		10		
2	3	1x4, 2x2, 1x2+2x1 and 4x1		12		
3	3	1x4, 2x2, 1x2+2x1 and 4x1		12		
	T	Intel_Hyper-Threading_Technology_	Intel_Virtualization_Technology_VTx_			
1	100 C	Yes	Yes			
2	100 C	Yes	Yes			
3	100 C	Yes	Yes			
	Intel_64_	Instruction_Set	Instruction_Set_Extensions		Idle_States	
1	Yes	64-bit	SSE4.1/4.2, AVX 2.0		Yes	
2	Yes	64-bit	SSE4.1/4.2, AVX 2.0		Yes	
3	Yes	64-bit	SSE4.1/4.2, AVX 2.0		Yes	

	Thermal_Monitoring_Technologies	Secure_Key	Execute_Disable_Bit
1	Yes	Yes	Yes
2	Yes	Yes	Yes
3	Yes	Yes	Yes

This command to import a data file of .csv form from the path into the program.

2.1.2 Import libraries

```

1 # Import libraries
2 library(dplyr)
3 library(stringr)
4 library(GGally)
5 library(corrplot)
6 library(caTools)
7 library(MASS)
8 library(car)
9 library(e1071)
10 library(nortest)

```

2.1.3 Explanation of the dataset

With the technological revolution, computers vary in purpose, leading to variation in their component, including variation of CPUs in statistical information. The dataset contains CPUs' vital information. All of them have a relationship with each other.

2.2 Data cleaning

2.2.1 Removing unused data

Checking how many missing values in each variable.

```

1 # Checking missing values
2 apply(is.na(CPU_data), 2, sum)

```

Product_Collection	Vertical_Segment
0	0
Processor_Number	Status
0	0
Launch_Date	Lithography
0	0
Recommended_Customer_Price	nb_of_Cores
0	0
nb_of_Threads	Processor_Base_Frequency
856	0
Max_Turbo_Frequency	Cache
0	0
Bus_Speed	TDP
0	0
Embedded_Options_Available	Conflict_Free
0	0
Max_Memory_Size	Memory_Types
0	0
Max_nb_of_Memory_Channels	Max_Memory_Bandwidth

	869		0
ECC_Memory_Supported		Processor_Graphics_	
	0		2283
Graphics_Base_Frequency		Graphics_Max_Dynamic_Frequency	
	0		0
Graphics_Video_Max_Memory		Graphics_Output	
	0		0
Support_4k		Max_Resolution_HDMI	
	2283		0
Max_Resolution_DP	Max_Resolution_eDP_Integrated_Flat_Panel		
	0		0
DirectX_Support		OpenGL_Support	
	0		2283
PCI_Express_Revision		PCI_Express_Configurations_	
	0		0
Max_nb_of_PCI_Express_Lanes		T	
	1104		0
Intel_Hyper_Threading_Technology_		Intel_Virtualization_Technology_VTx_	
	0		0
Intel_64_		Instruction_Set	
	0		0
Instruction_Set_Extensions		Idle_States	
	0		0
Thermal_Monitoring_Technologies		Secure_Key	
	0		0
Execute_Disable_Bit			
	0		

Values of **Processor_Graphics_**, **Support_4k**, and **OpenGL_Support** are missing in all rows; therefore, we remove columns of these variables.

Moreover, we will not use **Product_Collection**, **Vertical_Segment**, **Processor_Number**, **Status**, **Launch_Date**, **Instruction_Set_Extensions**, **PCI_Express_Configurations_**, **Max_Resolution_eDP_Integrated_Flat_Panel**, **Max_Resolution_DP**, **Graphics_Output**, **Max_Resolution_HDMI**, **Memory_Types**, **Bus_Speed**, **Cache**, **DirectX_Support**, **PCI_Express_Revision**, **Conflict_Free** to analyze, therefore, we also remove these variables.

```

1 # Remove unused variables
2 delete = c("OpenGL_Support", "Support_4k", "Processor_Graphics_", "
  Product_Collection", "Vertical_Segment", "Processor_Number", "Status",
  "Launch_Date", "Instruction_Set_Extensions", "PCI_Express_
  Configurations_", "Max_Resolution_eDP_Integrated_Flat_Panel", "Max_
  Resolution_DP", "Graphics_Output", "Max_Resolution_HDMI", "Memory_Types",
  "Bus_Speed", "Cache", "DirectX_Support", "PCI_Express_Revision", "
  Conflict_Free")
3
4 CPU_data <- CPU_data[, !names(CPU_data) %in% delete, drop = FALSE]

```

2.2.2 Filling missing values

We convert space and escape sequence to *NA* value.

```

1 # Convert " " and "\n- " to NA
2 CPU_data[(CPU_data == " ") | (CPU_data == "\n- ")] <- NA

```

We remove unit of some numerical variables and transform them to numeric form.

```
1 # Remove unit of numerical variables and transform to numeric form
2 CPU_data$Lithography <- as.numeric(sub("nm", "", CPU_data$Lithography))
3 CPU_data$Max_Turbo_Frequency <- as.numeric(sub("GHz", "", CPU_data$Max_
  Turbo_Frequency))
4 CPU_data$TDP <- as.numeric(sub("W", "", CPU_data$TDP))
5 CPU_data$Max_Memory_Size <- as.numeric(sub("GB", "", CPU_data$Max_
  Memory_Size))
6 CPU_data$Max_Memory_Bandwidth <- as.numeric(sub("GB/s", "", CPU_data$
  Max_Memory_Bandwidth))
7 CPU_data$Graphics_Base_Frequency <- as.numeric(sub("MHz", "", CPU_data$
  Graphics_Base_Frequency))
8 CPU_data$Graphics_Video_Max_Memory <- as.numeric(sub("GB", "", CPU_data
  $Graphics_Video_Max_Memory))
9 CPU_data$T <- as.numeric(sub("°C", "", CPU_data$T))
```

There are some numerical variables that need to be converted unit before removing units and transform to numeric form.

```
1 # Convert units for Graphics_Max_Dynamic_Frequency
2 subset_GHz <- CPU_data[grepl("GHz", CPU_data$Graphics_Max_Dynamic_
  Frequency, ignore.case = TRUE), ]
3 CPU_data <- CPU_data[!grepl("GHz", CPU_data$Graphics_Max_Dynamic_
  Frequency, ignore.case = TRUE), ]
4
5 subset_GHz$Graphics_Max_Dynamic_Frequency <- gsub("GHz", "", subset_GHz$
  Graphics_Max_Dynamic_Frequency, fixed = TRUE)
6 subset_GHz$Graphics_Max_Dynamic_Frequency <- as.numeric(subset_GHz$
  Graphics_Max_Dynamic_Frequency)
7 subset_GHz$Graphics_Max_Dynamic_Frequency <- subset_GHz$Graphics_Max_
  Dynamic_Frequency*(1000)
8
9 CPU_data$Graphics_Max_Dynamic_Frequency <- gsub("MHz", "", CPU_data$
  Graphics_Max_Dynamic_Frequency, fixed = TRUE)
10 CPU_data$Graphics_Max_Dynamic_Frequency <- as.numeric(CPU_data$Graphics
  _Max_Dynamic_Frequency)
11
12 CPU_data <- bind_rows(CPU_data, subset_GHz)
13
14 # Convert units for Processor_Base_Frequency
15 subset_GHz <- CPU_data[grepl("GHz", CPU_data$Processor_Base_Frequency,
  ignore.case = TRUE), ]
16 CPU_data <- CPU_data[!grepl("GHz", CPU_data$Processor_Base_Frequency,
  ignore.case = TRUE), ]
17
18 subset_GHz$Processor_Base_Frequency <- gsub("GHz", "", subset_GHz$
  Processor_Base_Frequency, fixed = TRUE)
19
20 subset_GHz$Processor_Base_Frequency <- as.numeric(subset_GHz$Processor_
```

```
Base_Frequency)
21 subset_GHz$Processor_Base_Frequency <- subset_GHz$Processor_Base_
    Frequency*(1000)
22
23 CPU_data$Processor_Base_Frequency <- gsub("MHz", "", CPU_data$Processor_
    Base_Frequency, fixed = TRUE)
24 CPU_data$Processor_Base_Frequency <- as.numeric(CPU_data$Processor_Base
    _Frequency)
25
26 CPU_data <- bind_rows(CPU_data, subset_GHz)
```

We fill missing values of numerical variables by their *mean* values.

```
1 # Filling missing values for numerical variables
2 CPU_data$Lithography[is.na(CPU_data$Lithography)] = mean(CPU_data$
    Lithography, na.rm=T)
3 CPU_data$nb_of_Cores[is.na(CPU_data$nb_of_Cores)] = mean(CPU_data$nb_of
    _Cores, na.rm=T)
4 CPU_data$nb_of_Threads[is.na(CPU_data$nb_of_Threads)] = mean(CPU_data$
    nb_of_Threads, na.rm=T)
5 CPU_data$Max_Turbo_Frequency[is.na(CPU_data$Max_Turbo_Frequency)] =
    mean(CPU_data$Max_Turbo_Frequency, na.rm=T)
6 CPU_data$TDP[is.na(CPU_data$TDP)] = mean(CPU_data$TDP, na.rm=T)
7 CPU_data$Max_Memory_Size[is.na(CPU_data$Max_Memory_Size)] = mean(CPU_
    data$Max_Memory_Size, na.rm=T)
8 CPU_data$Max_nb_of_Memory_Channels[is.na(CPU_data$Max_nb_of_Memory_
    Channels)] = mean(CPU_data$Max_nb_of_Memory_Channels, na.rm=T)
9 CPU_data$Max_Memory_Bandwidth[is.na(CPU_data$Max_Memory_Bandwidth)] =
    mean(CPU_data$Max_Memory_Bandwidth, na.rm=T)
10 CPU_data$Graphics_Base_Frequency[is.na(CPU_data$Graphics_Base_Frequency
    )] = mean(CPU_data$Graphics_Base_Frequency, na.rm=T)
11 CPU_data$Graphics_Video_Max_Memory[is.na(CPU_data$Graphics_Video_Max_
    Memory)] = mean(CPU_data$Graphics_Video_Max_Memory, na.rm=T)
12 CPU_data$Max_nb_of_PCI_Express_Lanes[is.na(CPU_data$Max_nb_of_PCI_
    Express_Lanes)] = mean(CPU_data$Max_nb_of_PCI_Express_Lanes, na.rm=T
    )
13 CPU_data$T[is.na(CPU_data$T)] = mean(CPU_data$T, na.rm=T)
14 CPU_data$Graphics_Max_Dynamic_Frequency[is.na(CPU_data$Graphics_Max_
    Dynamic_Frequency)] = mean(CPU_data$Graphics_Max_Dynamic_Frequency,
    na.rm=T)
15 CPU_data$Processor_Base_Frequency[is.na(CPU_data$Processor_Base_
    Frequency)] = mean(CPU_data$Processor_Base_Frequency, na.rm=T)
16
17 CPU_data$Recommended_Customer_Price <- gsub("$", "", CPU_data$
    Recommended_Customer_Price, fixed = TRUE)
18 CPU_data$Recommended_Customer_Price <- as.numeric(CPU_data$Recommended_
    Customer_Price)
```

We fill missing values of categorical variables by their *mode*.

```

1 # Filling missing values for categorical variables
2 fillmode <- function(column) {
3   mode_value <- names(sort(table(column), decreasing = TRUE))[1]
4   column[is.na(column)] <- mode_value
5   return(column)
6 }
7
8 fill_col <- c("Embedded_Options_Available", "ECC_Memory_Supported",
9             "Intel_Hyper_Threading_Technology_", "Intel_Virtualization_Technology_
10            VTx_", "Intel_64_", "Instruction_Set", "Idle_States", "Thermal_
11            Monitoring_Technologies", "Secure_Key", "Execute_Disable_Bit")
12
13 CPU_data[fill_col] <- lapply(CPU_data[fill_col], fillmode)

```

We transform categorical variables to numeric form by using categorical signs “1” and “0”. This step is conducted for training model.

```

1 # Transform categorical variables to numerical variables
2 CPU_data$Embedded_Options_Available <- ifelse(CPU_data$Embedded_Options
3         _Available == "Yes", 1, 0)
4 CPU_data$ECC_Memory_Supported <- ifelse(CPU_data$ECC_Memory_Supported
5         == "Yes", 1, 0)
6 CPU_data$Intel_Hyper_Threading_Technology_ <- ifelse(CPU_data$Intel_
7         Hyper_Threading_Technology_ == "Yes", 1, 0)
8 CPU_data$Intel_Virtualization_Technology_VTx_ <- ifelse(CPU_data$Intel_
9         Virtualization_Technology_VTx_ == "Yes", 1, 0)
10 CPU_data$Intel_64_ <- ifelse(CPU_data$Intel_64_ == "Yes", 1, 0)
11 CPU_data$Idle_States <- ifelse(CPU_data$Idle_States == "Yes", 1, 0)
12 CPU_data$Thermal_Monitoring_Technologies <- ifelse(CPU_data$Thermal_
13         Monitoring_Technologies == "Yes", 1, 0)
14 CPU_data$Secure_Key <- ifelse(CPU_data$Secure_Key == "Yes", 1, 0)
15 CPU_data$Execute_Disable_Bit <- ifelse(CPU_data$Execute_Disable_Bit ==
16         "Yes", 1, 0)
17
18 CPU_data$Instruction_Set <- gsub("32-bit", "0", CPU_data$Instruction_Set
19         , fixed = TRUE)
20 CPU_data$Instruction_Set <- gsub("64-bit", "1", CPU_data$Instruction_Set
21         , fixed = TRUE)
22 CPU_data$Instruction_Set <- gsub("Itanium 1", "2", CPU_data$Instruction_
23         Set, fixed = TRUE)
24 CPU_data$Instruction_Set <- as.numeric(CPU_data$Instruction_Set)

```

After all, we check the missing values again.

```

1 # Checking missing values
2 apply(is.na(CPU_data), 2, sum)

```

Lithography	Recommended_Customer_Price
0	1262
nb_of_Cores	Processor_Base_Frequency

	0	0
Max_Turbo_Frequency		TDP
	0	0
Embedded_Options_Available		Max_Memory_Size
	0	0
Max_Memory_Bandwidth		ECC_Memory_Supported
	0	0
Graphics_Base_Frequency		Graphics_Max_Dynamic_Frequency
	0	0
Graphics_Video_Max_Memory		T
	0	0
Intel_Hyper_Threading_Technology	Intel_Virtualization_Technology_VTx	
	0	0
Intel_64		Instruction_Set
	0	0
Idle_States	Thermal_Monitoring_Technologies	
	0	0
Secure_Key		Execute_Disable_Bit
	0	0

Since our objective of this report is to analyze and predict the recommended customer price of CPU, **Recommended_Customer_Price** is the only variable that has missing values after the pre-processing step.

2.3 Data visualization

2.3.1 Descriptive statistics for each variable

We divide CPU_data into 2 separate data CPU_learn and CPU_train, which CPU_train is the set that does not include missing values, even missing values in the Recommended_Customer_Price variable.

```

1 # Divide CPU_data into CPU_learn and CPU_train
2 CPU_learn <- subset(CPU_data, is.na(Recommended_Customer_Price))
3 CPU_train <- subset(CPU_data, !is.na(Recommended_Customer_Price))
4 summary(CPU_train)

```

Lithography	Recommended_Customer_Price	nb_of_Cores
Min. : 14.00	Min. : 2.54	Min. : 1.000
1st Qu.: 14.00	1st Qu.: 107.00	1st Qu.: 2.000
Median : 22.00	Median : 239.50	Median : 2.000
Mean : 26.57	Mean : 268.37	Mean : 3.148
3rd Qu.: 32.00	3rd Qu.: 378.00	3rd Qu.: 4.000
Max. : 130.00	Max. : 999.00	Max. : 16.000
Processor_Base_Frequency	Max_Turbo_Frequency	TDP
Min. : 32	Min. : 1.300	Min. : 0.025
1st Qu.: 1800	1st Qu.: 3.198	1st Qu.: 17.000
Median : 2300	Median : 3.198	Median : 35.000
Mean : 2304	Mean : 3.181	Mean : 44.061
3rd Qu.: 2800	3rd Qu.: 3.200	3rd Qu.: 60.242
Max. : 4100	Max. : 4.500	Max. : 140.000
Embedded_Options_Available	Max_Memory_Size	Max_Memory_Bandwidth
Min. : 0.0000	Min. : 1.00	Min. : 1.60
1st Qu.: 0.0000	1st Qu.: 16.00	1st Qu.: 25.60
Median : 0.0000	Median : 32.00	Median : 29.80
Mean : 0.3235	Mean : 78.03	Mean : 29.91
3rd Qu.: 1.0000	3rd Qu.: 128.00	3rd Qu.: 35.08
Max. : 1.0000	Max. : 768.00	Max. : 85.30

ECC_Memory_Supported	Graphics_Base_Frequency	Graphics_Max_Dynamic_Frequency
Min. :0.0000	Min. :100.0	Min. : 500
1st Qu.:0.0000	1st Qu.:327.5	1st Qu.:1000
Median :0.0000	Median :420.1	Median :1032
Mean :0.3914	Mean :405.9	Mean :1031
3rd Qu.:1.0000	3rd Qu.:420.1	3rd Qu.:1050
Max. :1.0000	Max. :900.0	Max. :1350
Graphics_Video_Max_Memory	T	Intel_Hyper_Threading_Technology
Min. : 1.00	Min. : 53.90	Min. :0.000
1st Qu.:22.81	1st Qu.: 73.30	1st Qu.:0.000
Median :22.81	Median : 83.25	Median :1.000
Mean :22.54	Mean : 86.18	Mean :0.621
3rd Qu.:22.81	3rd Qu.:100.00	3rd Qu.:1.000
Max. :64.00	Max. :110.00	Max. :1.000
Intel_Virtualization_Technology_VTx	Intel_64	Instruction_Set
Min. :0.000	Min. :0.0000	Min. :0.0000
1st Qu.:1.000	1st Qu.:1.0000	1st Qu.:1.0000
Median :1.000	Median :1.0000	Median :1.0000
Mean :0.942	Mean :0.9691	Mean :0.9679
3rd Qu.:1.000	3rd Qu.:1.0000	3rd Qu.:1.0000
Max. :1.000	Max. :1.0000	Max. :2.0000
Idle_States	Thermal_Monitoring_Technologies	Secure_Key
Min. :0.0000	Min. :0.0000	Min. :0.0000
1st Qu.:1.0000	1st Qu.:1.0000	1st Qu.:1.0000
Median :1.0000	Median :1.0000	Median :1.0000
Mean :0.9704	Mean :0.9259	Mean :0.9506
3rd Qu.:1.0000	3rd Qu.:1.0000	3rd Qu.:1.0000
Max. :1.0000	Max. :1.0000	Max. :1.0000
Execute_Disable_Bit		
Min. :0.0000		
1st Qu.:1.0000		
Median :1.0000		
Mean :0.9975		
3rd Qu.:1.0000		
Max. :1.0000		

2.3.2 Plotting graph

2.3.2.a Histogram

```


1 variables<-colnames(CPU_train)
2 # Histogram
3 not_his = c("Embedded_Options_Available", "ECC_Memory_Supported", "
  Intel_Hyper_Threading_Technology_", "Intel_Virtualization_Technology
  _VTx_", "Intel_64_", "Idle_States", "Thermal_Monitoring_Technologies
  ", "Secure_Key", "Execute_Disable_Bit", "Instruction_Set")
4 his_variables <- setdiff(variables, not_his)
5
6 for(var in his_variables){
7   temp <-gsub(",", "", CPU_train[var]);
8   temp <-as.numeric(unlist(CPU_train[var]));
9   hist(temp,
10     main=var,
11     col = "gray",

```


```

12   xlab="Units",
13   freq = FALSE,
14   cex.main = 1
15   );
16 }


```




Concentrate around a point but not symmetric




Symmetrically distributed




Exist an outlier




Right-skewed distribution




Concentrate around a point



Right-skewed distribution



Left-skewed distribution



Distribution not good

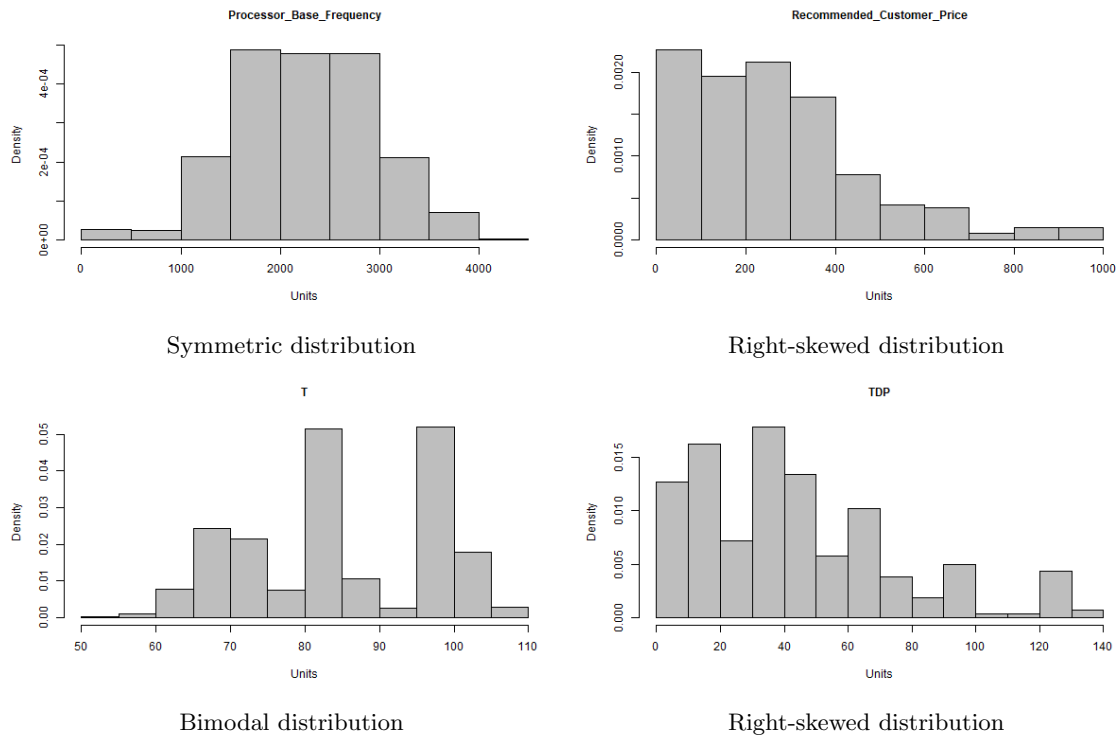





Figure 1: Histograms of data features

The data features have their own unique distribution, which contribute to the model later on.


2.3.2.b Box plot

```
1 for (i in variables){
2   boxplot(CPU_train[i],
3     col="orange",
4     xlab = i,
5     cex.lab = 1,
6     title.cex = 1,
7     border="brown")
8 }
```







Execute_Disable_Bit




Graphics_Base_Frequency




Graphics_Max_Dynamic_Frequency




Graphics_Video_Max_Memory




Graphics_Max_Dynamic_Frequency




Graphics_Video_Max_Memory




Idle_States




Instruction_Set




Intel_64_



Intel_Hyper-Threading_Technology_



Intel_Virtualization_Technology_VTx_



Lithography

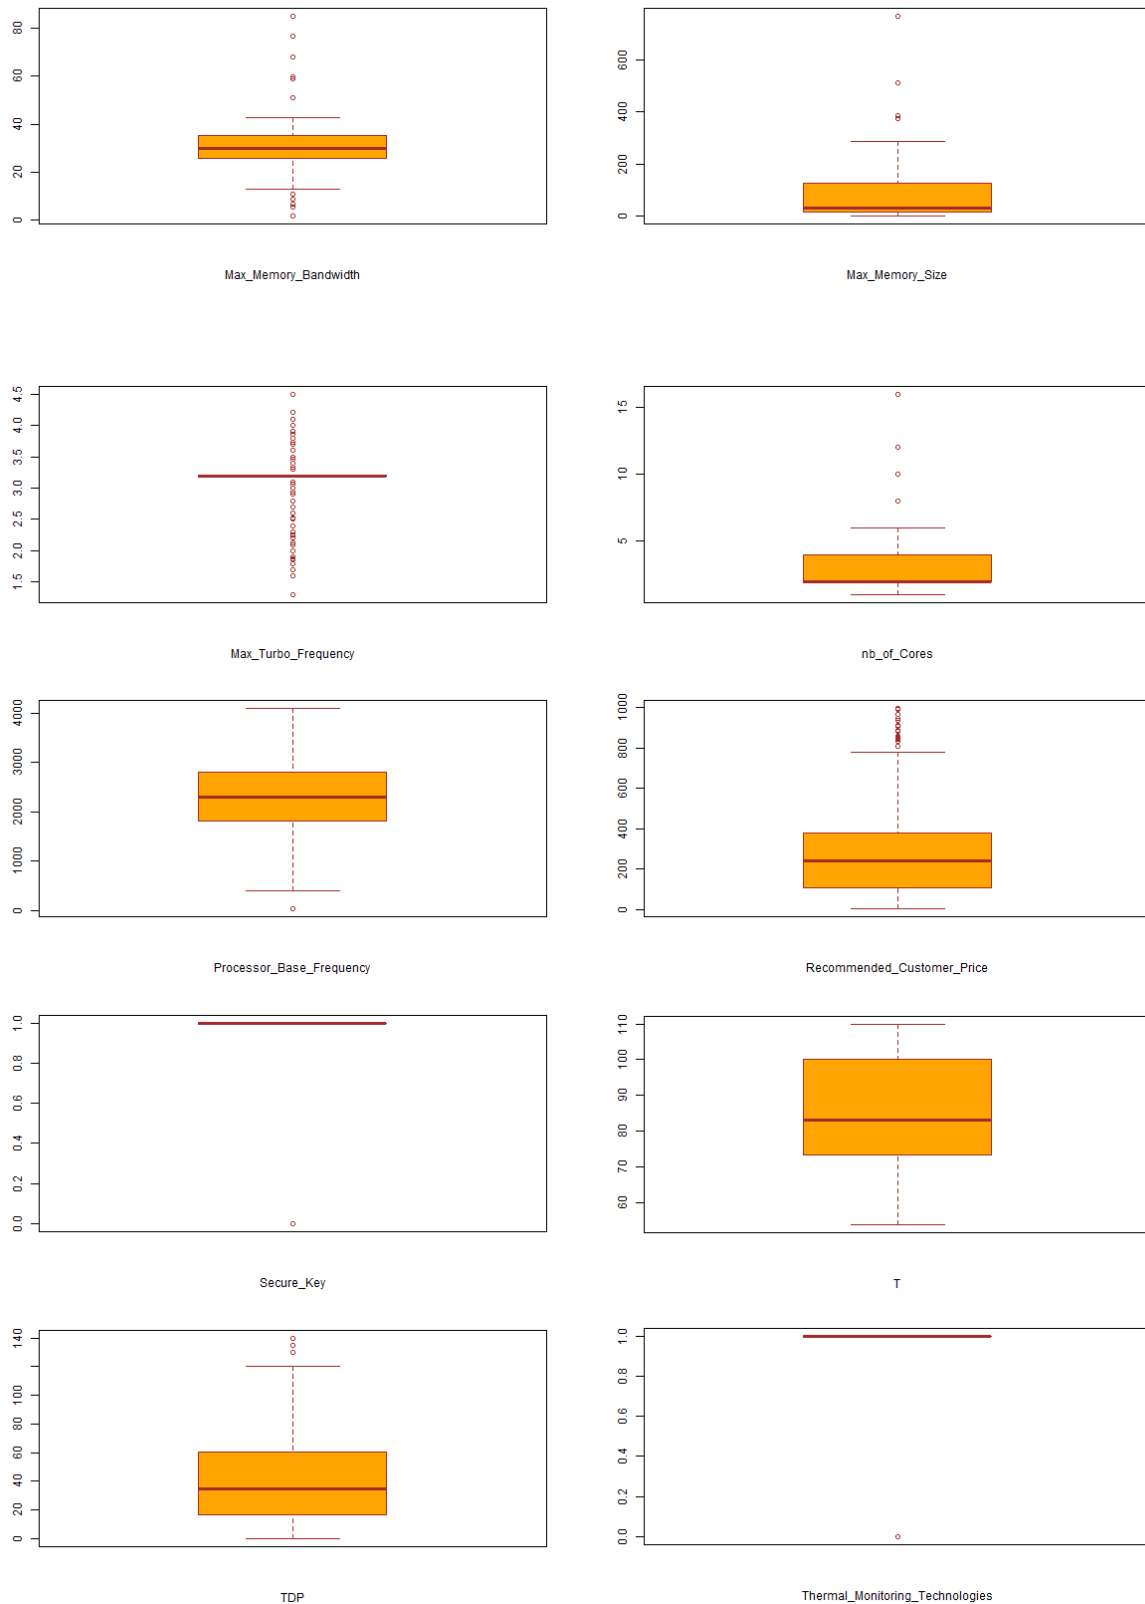


Figure 2: Scatter plots of data features

Some of the features are highly distributed and some other are highly concentrated.

2.3.2.c Correlation matrix

```

1 # Correlation
2 # Calculate the correlation matrix
3 select <- c("Processor_Base_Frequency", "Lithography", "nb_of_Cores", "
  Max_Turbo_Frequency", "Recommended_Customer_Price")
4 correlation_matrix <- cor(CPU_train[select])
5 # Plot the correlation matrix
6 corrplot(correlation_matrix, method = "color", type = "upper", order =
  "hclust", tl.col = "black", tl.srt = 45)

```




Figure 3: Correlation matrix (1)

```

1 # Calculate the correlation matrix
2 select <- c("Max_Memory_Bandwidth", "TDP", "Embedded_Options_Available
  ", "Max_Memory_Size", "Recommended_Customer_Price")
3 correlation_matrix <- cor(CPU_train[select])
4 # Plot the correlation matrix
5 corrplot(correlation_matrix, method = "color", type = "upper", order =
  "hclust", tl.col = "black", tl.srt = 45)

```

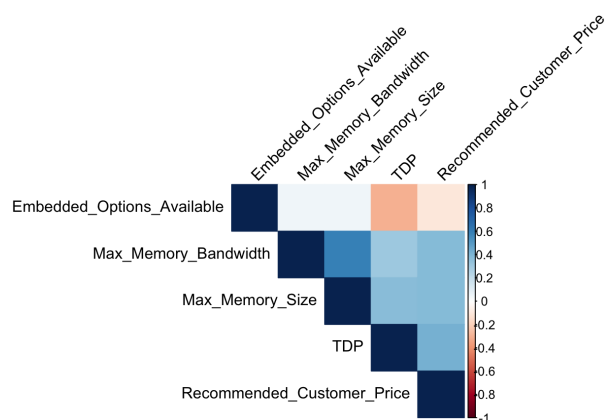


Figure 4: Correlation matrix (2)

```

1 # Calculate the correlation matrix

```

```

2 select <- c("ECC_Memory_Supported", "Graphics_Base_Frequency", "
  Graphics_Max_Dynamic_Frequency", "Graphics_Video_Max_Memory", "
  Recommended_Customer_Price")
3 correlation_matrix <- cor(CPU_train[select])
4
5 # Plot the correlation matrix
6 corrplot(correlation_matrix, method = "color", type = "upper", order =
  "hclust", tl.col = "black", tl.srt = 45)

```

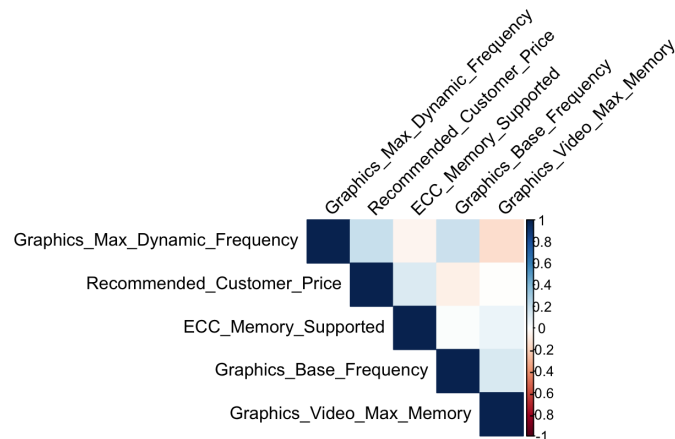


Figure 5: Correlation matrix (3)

```

1 # Calculate the correlation matrix
2 select <- c("T", "Intel_Hyper_Threading_Technology_", "Intel_
  Virtualization_Technology_VTx_" , "Intel_64_" , "Instruction_Set" , "
  Recommended_Customer_Price")
3 correlation_matrix <- cor(CPU_train[select])
4 # Plot the correlation matrix
5 corrplot(correlation_matrix, method = "color", type = "upper", order =
  "hclust", tl.col = "black", tl.srt = 45)

```




Figure 6: Correlation matrix (4)

```

1 # Calculate the correlation matrix
2 select <- c("Idle_States", "Thermal_Monitoring_Technologies", "
  Secure_Key", "Execute_Disable_Bit", "Recommended_Customer_Price")
3 correlation_matrix <- cor(CPU_train[select])
4 # Plot the correlation matrix
5 corplot(correlation_matrix, method = "color", type = "upper", order =
  "hclust", tl.col = "black", tl.srt = 45)

```




Figure 7: Correlation matrix (5)

In the correlation matrices, we respectively delete Lithography, Graphics Base Frequency, Graphics Video Max Memory, T, Idle states, Thermal monitoring technology due to low correlation.

2.4 Models building

2.4.1 Multivariate Linear Regression (MLR)

We divide CPU_train into 2 separate subsets, which are the train set and the test set. This is good practice for training prediction model.

```

1 # Divide CPU_train into train_df and test_df
2 set.seed(42)
3 # Use 70% of dataset as training set and 30% as test set
4 split <- sample.split(CPU_train, SplitRatio = 0.70)
5 train_df <- subset(CPU_train, split == TRUE)
6 test_df <- subset(CPU_train, split == FALSE)
7 # Fitting model
8 model1 <- lm(Recommended_Customer_Price ~ ., data = train_df)
9 summary(model1)

```

```

## Call:
## lm(formula = Recommended_Customer_Price ~ ., data = train_df)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -331.63  -62.02   -2.56   49.79  689.49
##
## Coefficients:

```



```
##               Estimate Std. Error t value Pr(>|t|)
## (Intercept)    -9.696e+02  2.311e+02  -4.195 3.19e-05 ***
## Lithography     2.607e+00  6.131e-01   4.253 2.49e-05 ***
## nb_of_Cores     3.952e+01  4.101e+00   9.636 < 2e-16 ***
## Processor_Base_Frequency -8.799e-03  1.362e-02  -0.646 0.518484
## Max_Turbo_Frequency 3.212e+01  1.783e+01   1.802 0.072143 .
## TDP            1.542e+00  3.500e-01   4.405 1.28e-05 ***
## Embedded_Options_Available -1.626e+01  1.497e+01  -1.087 0.277737
## Max_Memory_Size  2.732e-01  7.160e-02   3.816 0.000151 ***
## Max_Memory_Bandwidth 1.999e+00  7.553e-01   2.647 0.008371 **
## ECC_Memory_Supported 4.462e+00  1.367e+01   0.326 0.744315
## Graphics_Base_Frequency -7.534e-02  4.498e-02  -1.675 0.094523 .
## Graphics_Max_Dynamic_Frequency 1.845e-01  6.260e-02   2.947 0.003351 **
## Graphics_Video_Max_Memory -7.518e-01  3.912e-01  -1.922 0.055146 .
## T              3.721e+00  5.665e-01   6.569 1.21e-10 ***
## Intel_Hyper_Threading_Technology_ 1.331e+02  1.217e+01  10.937 < 2e-16 ***
## Intel_Virtualization_Technology_VTx_ 1.004e+02  3.127e+01   3.210 0.001406 **
## Intel_64_       -1.778e+02  8.299e+01  -2.143 0.032589 *
## Instruction_Set  2.623e+02  7.430e+01   3.530 0.000452 ***
## Idle_States     -5.362e+01  3.335e+01  -1.608 0.108501
## Thermal_Monitoring_Technologies 1.815e+01  2.669e+01   0.680 0.496678
## Secure_Key       1.313e+01  3.024e+01   0.434 0.664268
## Execute_Disable_Bit 1.141e+02  2.020e+02   0.565 0.572461
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 124.5 on 530 degrees of freedom
## Multiple R-squared:  0.6135, Adjusted R-squared:  0.5982
## F-statistic: 40.07 on 21 and 530 DF, p-value: < 2.2e-16
```

We are going to use stepwise regression to minimize the number of predictors.

```
1 # Stepwise
2 model2 <- stepAIC(model1, direction = "both")
3 summary(model2)

## Call:
## lm(formula = Recommended_Customer_Price ~ Lithography + nb_of_Cores +
##     Max_Turbo_Frequency + TDP + Max_Memory_Size + Max_Memory_Bandwidth +
##     Graphics_Base_Frequency + Graphics_Max_Dynamic_Frequency +
##     Graphics_Video_Max_Memory + T + Intel_Hyper_Threading_Technology_ +
##     Intel_Virtualization_Technology_VTx_ + Intel_64_ + Instruction_Set +
##     Idle_States, data = train_df)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -328.09  -62.43   -1.53   55.45   691.04
##
## Coefficients:
##               Estimate Std. Error t value Pr(>|t|)
## (Intercept)    -853.83393  113.93898  -7.494 2.78e-13 ***
## Lithography      2.56372    0.55937   4.583 5.70e-06 ***
## nb_of_Cores     39.76142    3.89737  10.202 < 2e-16 ***
## Max_Turbo_Frequency 32.48697  16.57622   1.960 0.050531 .
## TDP            1.44426    0.28265   5.110 4.49e-07 ***
## Max_Memory_Size  0.27880    0.06897   4.043 6.06e-05 ***
## Max_Memory_Bandwidth 2.04415    0.71365   2.864 0.004342 **
## Graphics_Base_Frequency -0.08689    0.04274  -2.033 0.042551 *
## Graphics_Max_Dynamic_Frequency 0.17887    0.05798   3.085 0.002141 **
```

```
## Graphics_Video_Max_Memory      -0.84533    0.37893   -2.231  0.026103 *
## T                               3.65881    0.54728    6.685  5.79e-11 ***
## Intel_Hyper_Threading_Technology_ 133.41953   11.70402   11.399 < 2e-16 ***
## Intel_Virtualization_Technology_VTx_ 108.74411   28.33553    3.838  0.000139 ***
## Intel_64_                      -145.95678   53.21008   -2.743  0.006291 **
## Instruction_Set                 238.27994   49.66475    4.798  2.08e-06 ***
## Idle_States                    -45.87275   32.49853   -1.412  0.158668
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 124.1 on 536 degrees of freedom
## Multiple R-squared:  0.6118, Adjusted R-squared:  0.6009
## F-statistic:  56.3 on 15 and 536 DF,  p-value: < 2.2e-16
```

2.4.2 Analysis of Variance (ANOVA)

```
1 multiAnova <- aov(Recommended_Customer_Price ~ Lithography + nb_of_
  Cores + Max_Turbo_Frequency + TDP + Max_Memory_Size + Max_Memory_
  Bandwidth + Graphics_Base_Frequency + Graphics_Max_Dynamic_
  Frequency + Graphics_Video_Max_Memory + T + Intel_Hyper_Threading_
  Technology_ + Intel_Virtualization_Technology_VTx_ + Intel_64_ +
  Instruction_Set + Idle_States, data = train_df)
2 summary(multiAnova)
```

```
##              Df Sum Sq Mean Sq F value    Pr(>F)
## Lithography      1    1231      1231    0.080 0.777447
## nb_of_Cores      1 5560077 5560077 361.297 < 2e-16 ***
## Max_Turbo_Frequency 1  894154   894154   58.103 1.14e-13 ***
## TDP              1 1020681 1020681   66.324 2.70e-15 ***
## Max_Memory_Size  1  679416   679416   44.149 7.49e-11 ***
## Max_Memory_Bandwidth 1  127876   127876    8.309 0.004102 **
## Graphics_Base_Frequency 1  295168   295168   19.180 1.43e-05 ***
## Graphics_Max_Dynamic_Frequency 1  173493   173493   11.274 0.000842 ***
## Graphics_Video_Max_Memory 1  105013   105013    6.824 0.009247 **
## T                1 1288664 1288664   83.738 < 2e-16 ***
## Intel_Hyper_Threading_Technology_ 1 1950353 1950353 126.735 < 2e-16 ***
## Intel_Virtualization_Technology_VTx_ 1  407478   407478   26.478 3.75e-07 ***
## Intel_64_        1   49410    49410    3.211 0.073722 .
## Instruction_Set   1  413499   413499   26.869 3.09e-07 ***
## Idle_States      1   30662    30662    1.992 0.158668
## Residuals       536 8248616   15389
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
1 av_residual= rstandard(aov(Recommended_Customer_Price~.,
2                       data =CPU_train))
3 shapiro.test(av_residual)
```

```
## Shapiro-Wilk normality test
##
## data:  av_residual
## W = 0.94051, p-value < 2.2e-16
```

This dataset does not have normal distribution since $p\text{-value} < 0.05$, it means MLR is not efficient enough. Therefore, we apply the support vector machine method.

2.4.3 Support vector machine (SVM)

```
1 model3 <- svm(Recommended_Customer_Price ~ Lithography + nb_of_Cores +
  Max_Turbo_Frequency + TDP + Max_Memory_Size + Max_Memory_Bandwidth +
  Graphics_Base_Frequency + Graphics_Max_Dynamic_Frequency +
  Graphics_Video_Max_Memory + T + Intel_Hyper_Threading_Technology_ +
  Intel_Virtualization_Technology_VTx_ + Intel_64_ + Instruction_Set +
  Idle_States, data = train_df)
```

2.5 Model comparison

```
1 anova(model1, model2)

## Analysis of Variance Table
##
## Model 1: Recommended_Customer_Price ~ Lithography + nb_of_Cores +
##   Processor_Base_Frequency +
##     Max_Turbo_Frequency + TDP + Embedded_Options_Available +
##     Max_Memory_Size + Max_Memory_Bandwidth + ECC_Memory_Supported +
##     Graphics_Base_Frequency + Graphics_Max_Dynamic_Frequency +
##     Graphics_Video_Max_Memory + T + Intel_Hyper_Threading_Technology_ +
##     Intel_Virtualization_Technology_VTx_ + Intel_64_ + Instruction_Set +
##     Idle_States + Thermal_Monitoring_Technologies + Secure_Key +
##     Execute_Disable_Bit
## Model 2: Recommended_Customer_Price ~ Lithography + nb_of_Cores + Max_Turbo_Frequency
##   +
##     TDP + Max_Memory_Size + Max_Memory_Bandwidth + Graphics_Base_Frequency +
##     Graphics_Max_Dynamic_Frequency + Graphics_Video_Max_Memory +
##     T + Intel_Hyper_Threading_Technology_ + Intel_Virtualization_Technology_VTx_ +
##     Intel_64_ + Instruction_Set + Idle_States
##   Res.Df      RSS Df Sum of Sq      F Pr(>F)
## 1      530 8210950
## 2      536 8248616  -6      -37665  0.4052  0.8757
```

With p-value greater than 0.05, model1 and model2 have equivalent efficiency. However, R-squared of model2 is greater than model1 $0.6009 > 0.5982$, model2 is a little better than model1. Therefore, we will compare accuracy between model2 and model3.

```
1 pred_test=predict(model1,newdata=test_df)
2 SSE <- sum((test_df$Recommended_Customer_Price - pred_test) ^ 2) #Sum
  of Squares Error
3 SST <- sum((test_df$Recommended_Customer_Price - mean(test_df$
  Recommended_Customer_Price)) ^ 2) #Sum of Squares Total
4 cat("The accuracy of the model on test dataset: ",round((1 - SSE/SST)*
  100,2),"%")
5
6 pred_test=predict(model2,newdata=test_df)
7 SSE <- sum((test_df$Recommended_Customer_Price - pred_test) ^ 2) #Sum
  of Squares Error
8 SST <- sum((test_df$Recommended_Customer_Price - mean(test_df$
  Recommended_Customer_Price)) ^ 2) #Sum of Squares Total
9 cat("The accuracy of the model on test dataset: ",round((1 - SSE/SST)*
  100,2),"%")
```



```
## The accuracy of the model on test dataset: 60.32 %  
## The accuracy of the model on test dataset: 62.28 %
```

Accuracy of model3 is 62.28%, which is more than accuracy of model2 (60.32%), thus, model3 is the final model.

3 Conclusion

In this report, we imported the 'Intel_CPUs.csv' file for examination and performed data cleaning and data visualizations. Subsequently, we analyzed the processed dataset of 2283 CPU observations with 22 variables, including one target variable, "Recommended_Customer_Price", which contained missing values. To address these missing values, we employed linear models, evaluating the performance of various models and selecting the best fit model as our final method.

Our findings demonstrate the effectiveness of linear models in imputing missing values for CPU prices. The selected model successfully predicted missing values with high accuracy, utilizing the information provided by the remaining 21 variables. The analysis of the dataset also revealed several insights into the factors influencing CPU prices. These findings highlight the importance of these hardware specifications in determining the overall value of a CPU. This approach allowed us to utilize the complete dataset for further analysis and modeling efforts.

Despite our best efforts, due to the lack of knowledge and techniques and the limitations of the dataset, there might be some inaccuracies in this report. We would greatly appreciate your valuable feedback to help us improve our analysis and make it more robust in the future.

4 R script

```
1 ## Import data
2 # Import file
3 CPU_data = read.csv("~/Desktop/XSTK Assignment/work/Intel_CPUs.csv")
4 head(CPU_data, 3)
5 # Import libraries
6 library(dplyr)
7 library(stringr)
8 library(GGally)
9 library(corrplot)
10 library(caTools)
11 library(MASS)
12 library(car)
13 library(e1071)
14 library(nortest)
15 ## Data cleaning
16 # Checking missing values
17 apply(is.na(CPU_data), 2, sum)
18 # Remove unused variables
19 delete = c("OpenGL_Support", "Support_4k", "Processor_Graphics_", "
    Product_Collection", "Vertical_Segment", "Processor_Number", "Status
    ", "Launch_Date", "nb_of_Threads", "Instruction_Set_Extensions", "PCI_
    Express_Configurations_", "Max_Resolution_eDP_Integrated_Flat_Panel",
    "Max_Resolution_DP", "Graphics_Output", "Max_Resolution_HDMI", "Memory_
    Types", "Bus_Speed", "Cache", "DirectX_Support", "Max_nb_of_PCI_Express_
    Lanes", "Max_nb_of_Memory_Channels", "PCI_Express_Revision", "
    Conflict_Free")
20 CPU_data <- CPU_data[, !names(CPU_data) %in% delete, drop = FALSE]
21 # Convert " " and "\n- " to NA
22 CPU_data[(CPU_data == " ") | (CPU_data == "\n- ")] <- NA
23 # Remove unit of numerical variables and transform to numeric form
24 CPU_data$Lithography <- as.numeric(sub("nm", "", CPU_data$Lithography))
25 CPU_data$Max_Turbo_Frequency <- as.numeric(sub("GHz", "", CPU_data$Max_
    Turbo_Frequency))
26 CPU_data$TDP <- as.numeric(sub("W", "", CPU_data$TDP))
27 CPU_data$Max_Memory_Size <- as.numeric(sub("GB", "", CPU_data$Max_
    Memory_Size))
28 CPU_data$Max_Memory_Bandwidth <- as.numeric(sub("GB/s", "", CPU_data$
    Max_Memory_Bandwidth))
29 CPU_data$Graphics_Base_Frequency <- as.numeric(sub("MHz", "", CPU_data$
    Graphics_Base_Frequency))
30 CPU_data$Graphics_Video_Max_Memory <- as.numeric(sub("GB", "", CPU_data
    $Graphics_Video_Max_Memory))
31 CPU_data$T <- as.numeric(sub(" C ", "", CPU_data$T))
32 # Convert units for Graphics_Max_Dynamic_Frequency
33 subset_GHz <- CPU_data[grepl("GHz", CPU_data$Graphics_Max_Dynamic_
```

```
Frequency, ignore.case = TRUE) , ]
34 CPU_data <- CPU_data[!grepl("GHz", CPU_data$Graphics_Max_Dynamic_
    Frequency, ignore.case = TRUE), ]
35 subset_GHz$Graphics_Max_Dynamic_Frequency <- gsub("GHz", "", subset_GHz$
    Graphics_Max_Dynamic_Frequency, fixed = TRUE)
36 subset_GHz$Graphics_Max_Dynamic_Frequency <- as.numeric(subset_GHz$
    Graphics_Max_Dynamic_Frequency)
37 subset_GHz$Graphics_Max_Dynamic_Frequency <- subset_GHz$Graphics_Max_
    Dynamic_Frequency*(1000)
38 CPU_data$Graphics_Max_Dynamic_Frequency <- gsub("MHz", "", CPU_data$
    Graphics_Max_Dynamic_Frequency, fixed = TRUE)
39 CPU_data$Graphics_Max_Dynamic_Frequency <- as.numeric(CPU_data$Graphics
    _Max_Dynamic_Frequency)
40 CPU_data <- bind_rows(CPU_data, subset_GHz)
41 # Convert units for Processor_Base_Frequency
42 subset_GHz <- CPU_data[grepl("GHz", CPU_data$Processor_Base_Frequency,
    ignore.case = TRUE) , ]
43 CPU_data <- CPU_data[!grepl("GHz", CPU_data$Processor_Base_Frequency,
    ignore.case = TRUE), ]
44 subset_GHz$Processor_Base_Frequency <- gsub("GHz", "", subset_GHz$
    Processor_Base_Frequency, fixed = TRUE)
45 subset_GHz$Processor_Base_Frequency <- as.numeric(subset_GHz$Processor_
    Base_Frequency)
46 subset_GHz$Processor_Base_Frequency <- subset_GHz$Processor_Base_
    Frequency*(1000)
47 CPU_data$Processor_Base_Frequency <- gsub("MHz", "", CPU_data$Processor_
    Base_Frequency, fixed = TRUE)
48 CPU_data$Processor_Base_Frequency <- as.numeric(CPU_data$Processor_Base
    _Frequency)
49 CPU_data <- bind_rows(CPU_data, subset_GHz)
50 # Filling missing values for numerical variables
51 CPU_data$Lithography[is.na(CPU_data$Lithography)] = mean(CPU_data$
    Lithography, na.rm=T)
52 CPU_data$nb_of_Cores[is.na(CPU_data$nb_of_Cores)] = mean(CPU_data$nb_of
    _Cores, na.rm=T)
53 CPU_data$Max_Turbo_Frequency[is.na(CPU_data$Max_Turbo_Frequency)] =
    mean(CPU_data$Max_Turbo_Frequency, na.rm=T)
54 CPU_data$TDP[is.na(CPU_data$TDP)] = mean(CPU_data$TDP, na.rm=T)
55 CPU_data$Max_Memory_Size[is.na(CPU_data$Max_Memory_Size)] = mean(CPU_
    data$Max_Memory_Size, na.rm=T)
56 CPU_data$Max_Memory_Bandwidth[is.na(CPU_data$Max_Memory_Bandwidth)] =
    mean(CPU_data$Max_Memory_Bandwidth, na.rm=T)
57 CPU_data$Graphics_Base_Frequency[is.na(CPU_data$Graphics_Base_Frequency
    )] = mean(CPU_data$Graphics_Base_Frequency, na.rm=T)
58 CPU_data$Graphics_Video_Max_Memory[is.na(CPU_data$Graphics_Video_Max_
    Memory)] = mean(CPU_data$Graphics_Video_Max_Memory, na.rm=T)
59 CPU_data$T[is.na(CPU_data$T)] = mean(CPU_data$T, na.rm=T)
```

```
60 CPU_data$Graphics_Max_Dynamic_Frequency[is.na(CPU_data$Graphics_Max_
    Dynamic_Frequency)] = mean(CPU_data$Graphics_Max_Dynamic_Frequency,
    na.rm=T)
61 CPU_data$Processor_Base_Frequency[is.na(CPU_data$Processor_Base_
    Frequency)] = mean(CPU_data$Processor_Base_Frequency, na.rm=T)
62 CPU_data$Recommended_Customer_Price <- gsub("$", "", CPU_data$
    Recommended_Customer_Price, fixed = TRUE)
63 CPU_data$Recommended_Customer_Price <- as.numeric(CPU_data$Recommended_
    Customer_Price)
64 # Filling missing values for categorical variables
65 fillmode <- function(column) {
66     mode_value <- names(sort(table(column), decreasing = TRUE))[1]
67     column[is.na(column)] <- mode_value
68     return(column)
69 }
70 fill_col <- c("Embedded_Options_Available", "ECC_Memory_Supported", "
    Intel_Hyper_Threading_Technology_", "Intel_Virtualization_Technology_
    VTx_", "Intel_64_", "Instruction_Set", "Idle_States", "Thermal_
    Monitoring_Technologies", "Secure_Key", "Execute_Disable_Bit")
71 CPU_data[fill_col] <- lapply(CPU_data[fill_col], fillmode)
72 # Transform categorical variables to numerical variables
73 CPU_data$Embedded_Options_Available <- ifelse(CPU_data$Embedded_Options
    _Available == "Yes", 1, 0)
74 CPU_data$ECC_Memory_Supported <- ifelse(CPU_data$ECC_Memory_Supported
    == "Yes", 1, 0)
75 CPU_data$Intel_Hyper_Threading_Technology_ <- ifelse(CPU_data$Intel_
    Hyper_Threading_Technology_ == "Yes", 1, 0)
76 CPU_data$Intel_Virtualization_Technology_VTx_ <- ifelse(CPU_data$Intel_
    Virtualization_Technology_VTx_ == "Yes", 1, 0)
77 CPU_data$Intel_64_ <- ifelse(CPU_data$Intel_64_ == "Yes", 1, 0)
78 CPU_data$Idle_States <- ifelse(CPU_data$Idle_States == "Yes", 1, 0)
79 CPU_data$Thermal_Monitoring_Technologies <- ifelse(CPU_data$Thermal_
    Monitoring_Technologies == "Yes", 1, 0)
80 CPU_data$Secure_Key <- ifelse(CPU_data$Secure_Key == "Yes", 1, 0)
81 CPU_data$Execute_Disable_Bit <- ifelse(CPU_data$Execute_Disable_Bit ==
    "Yes", 1, 0)
82 CPU_data$Instruction_Set <- gsub("32-bit", "0", CPU_data$Instruction_Set
    ,fixed = TRUE)
83 CPU_data$Instruction_Set <- gsub("64-bit", "1", CPU_data$Instruction_Set
    ,fixed = TRUE)
84 CPU_data$Instruction_Set <- gsub("Itanium 1", "2", CPU_data$Instruction_
    Set,fixed = TRUE)
85 CPU_data$Instruction_Set <- as.numeric(CPU_data$Instruction_Set)
86 # Checking missing values
87 apply(is.na(CPU_data), 2, sum)
88
89 ## Data visualization
```



```
90 # Divide data set
91 CPU_learn <- subset(CPU_data, is.na(Recommended_Customer_Price))
92 CPU_train <- subset(CPU_data, !is.na(Recommended_Customer_Price))
93 summary(CPU_train)
94 # Variable for plotting
95 variables<-colnames(CPU_train)
96 # Histogram
97 not_his = c("Embedded_Options_Available", "ECC_Memory_Supported", "
          Intel_Hyper_Threading_Technology_", "Intel_Virtualization_Technology
          _VTx_", "Intel_64_", "Idle_States", "Thermal_Monitoring_Technologies
          ", "Secure_Key", "Execute_Disable_Bit", "Instruction_Set")
98 his_variables <- setdiff(variables, not_his)
99 for(var in his_variables){
100   temp <-gsub(",", "", CPU_train[var]);
101   temp <-as.numeric(unlist(CPU_train[var]));
102   hist(temp,
103         main=var,
104         col = "gray",
105         xlab="Units",
106         freq = FALSE,
107         cex.main = 1
108   );
109 }
110 # Box plot
111 for (i in variables){
112   boxplot(CPU_train[i],
113           col="orange",
114           xlab = i,
115           cex.lab = 1,
116           title.cex = 1,
117           border="brown")
118 }
119 # Correlation
120 # Calculate the correlation matrix
121 select <- c("Processor_Base_Frequency", "Lithography","nb_of_Cores", "
          Max_Turbo_Frequency", "Recommended_Customer_Price")
122 correlation_matrix <- cor(CPU_train[select])
123 # Plot the correlation matrix
124 corrplot(correlation_matrix, method = "color", type = "upper", order =
          "hclust", tl.col = "black", tl.srt = 45)
125 # Calculate the correlation matrix
126 select <- c("Max_Memory_Bandwidth", "TDP", "Embedded_Options_Available
          ", "Max_Memory_Size", "Recommended_Customer_Price")
127 correlation_matrix <- cor(CPU_train[select])
128 # Plot the correlation matrix
129 corrplot(correlation_matrix, method = "color", type = "upper", order =
          "hclust", tl.col = "black", tl.srt = 45)
```

```
130 # Calculate the correlation matrix
131 select <- c("ECC_Memory_Supported", "Graphics_Base_Frequency", "
Graphics_Max_Dynamic_Frequency", "Graphics_Video_Max_Memory", "
Recommended_Customer_Price")
132 correlation_matrix <- cor(CPU_train[select])
133 # Plot the correlation matrix
134 corrplot(correlation_matrix, method = "color", type = "upper", order =
"hc1ust", tl.col = "black", tl.srt = 45)
135 # Calculate the correlation matrix
136 select <- c("T", "Intel_Hyper_Threading_Technology_", "Intel_
Virtualization_Technology_VTx_" , "Intel_64_" , "Instruction_Set" , "
Recommended_Customer_Price")
137 correlation_matrix <- cor(CPU_train[select])
138 # Plot the correlation matrix
139 corrplot(correlation_matrix, method = "color", type = "upper", order =
"hc1ust", tl.col = "black", tl.srt = 45)
140 # Calculate the correlation matrix
141 select <- c("Idle_States", "Thermal_Monitoring_Technologies", "
Secure_Key" , "Execute_Disable_Bit" , "Recommended_Customer_Price")
142 correlation_matrix <- cor(CPU_train[select])
143 # Plot the correlation matrix
144 corrplot(correlation_matrix, method = "color", type = "upper", order =
"hc1ust", tl.col = "black", tl.srt = 45)
145
146 ## Models building
147 # Divide CPU_train into train_df and test_df
148 set.seed(42)
149 # Use 70% of dataset as training set and 30% as test set
150 split <- sample.split(CPU_train, SplitRatio = 0.70)
151 train_df <- subset(CPU_train, split == TRUE)
152 test_df <- subset(CPU_train, split == FALSE)
153 # Fitting model
154 model1 <- lm(Recommended_Customer_Price ~ ., data = train_df)
155 summary(model1)
156 # Stepwise regression
157 model2 <- stepAIC(model1, direction = "both")
158 summary(model2)
159 # Multi-factor ANOVA
160 multiAnova <- aov(Recommended_Customer_Price ~ Lithography + nb_of_
Cores + Max_Turbo_Frequency + TDP + Max_Memory_Size + Max_Memory_
Bandwidth + Graphics_Base_Frequency + Graphics_Max_Dynamic_
Frequency + Graphics_Video_Max_Memory + T + Intel_Hyper_Threading_
Technology_ + Intel_Virtualization_Technology_VTx_ + Intel_64_ +
Instruction_Set + Idle_States, data = train_df)
161 summary(multiAnova)
162 # Apply support vector machine
163 model3 <- svm(Recommended_Customer_Price ~ Lithography + nb_of_Cores +
```

```
Max_Turbo_Frequency + TDP + Max_Memory_Size + Max_Memory_Bandwidth +  
  Graphics_Base_Frequency + Graphics_Max_Dynamic_Frequency +  
  Graphics_Video_Max_Memory + T + Intel_Hyper_Threading_Technology_ +  
  Intel_Virtualization_Technology_VTx_ + Intel_64_ + Instruction_Set +  
  Idle_States, data = train_df)  
164  
165 ## Models comparison  
166 anova(model1, model2)  
167 pred_test = predict(model2, newdata = test_df)  
168 SSE <- sum((test_df$Recommended_Customer_Price - pred_test) ^ 2) #Sum  
  of Squares Error  
169 SST <- sum((test_df$Recommended_Customer_Price - mean(test_df$  
  Recommended_Customer_Price)) ^ 2) #Sum of Squares Total  
170 cat("The accuracy of the model on test dataset: ",round((1 - SSE/SST)*  
  100,2),"%")  
171 pred_test = predict(model3, newdata = test_df)  
172 SSE <- sum((test_df$Recommended_Customer_Price - pred_test) ^ 2) #Sum  
  of Squares Error  
173 SST <- sum((test_df$Recommended_Customer_Price - mean(test_df$  
  Recommended_Customer_Price)) ^ 2) #Sum of Squares Total  
174 cat("The accuracy of the model on test dataset: ",round((1 - SSE/SST)*  
  100,2),"%")
```

References

- [1] Douglas C. Montgomery, George C. Runger, *Applied Statistics and Probability for Engineers*, Wiley, Hoboken, NJ, 2019.
- [2] Prof. Nabendu Pal (2023), *Probability and Statistics lectures' notes*.
- [3] GeeksforGeeks, *Support Vector Machine (SVM) algorithm*. Available at: <https://www.geeksforgeeks.org/support-vector-machine-algorithm/>